

WI5118: AM Internship (cohort 2023)

Line Up Automation

Philip Soliman (4945255)

November 27, 2024



University supervisor:

Dr. Ir. L. Rens
Delft University of Technology
Faculty of Electrical Engineering, Mathematics and
Computer Science

Company supervisor:
J. Ramlakhan
Senior R&D Engineer
Research & Development

Contents

1	Company description	3
2	Problem description	5
2.1	Main requirements for automated line-up	5
2.2	Current situation	5
2.3	Internship goals	6
3	Approach	7
3.1	Understand automated line-up	7
3.2	Familiarize with the code base	7
3.3	Create a profiler for the code	7
3.4	Optimize the code	7
3.5	Identify unreadable code	7
3.6	Determine and improve upon current accuracy	8
3.7	Expand the fitting model to recognise ILUC	8
4	Understanding line up automation	9
4.1	The U-frame	9
4.2	Laser line scanners	9
4.3	Order of operations automated line-up	11
4.4	Flow of line up automation software	13
5	Profiler	15
5.1	Profiling test setup	15
5.2	Profiling results	17
6	Optimize & refactor code	19
7	Accuracy of circle fits	21
7.1	Statistical model	21
7.2	Accuracy estimation	21
8	ILUC Recognition	23
9	Conclusion	24
10	Skills learned	25

1 Company description

Allseas Engineering B.V. is a Dutch company that specializes in offshore pipeline installation and subsea construction. The company was founded in 1985 by Edward Heerema and is headquartered in Delft, the Netherlands. Allseas is known for its innovative engineering solutions and state-of-the-art vessels, such as the Pioneering Spirit, which is the largest construction vessel in the world. The company has a global presence with offices in the Netherlands, Switzerland, and the United Kingdom, and employs over 2,500 people worldwide.

Allseas' core business is the installation of subsea pipelines and the construction of offshore structures for the oil and gas industry. The company's vessels are equipped with advanced technology and equipment to perform a wide range of offshore operations. [1]

Allseas' organizational structure is divided into several departments, which are shown in Figure 1.

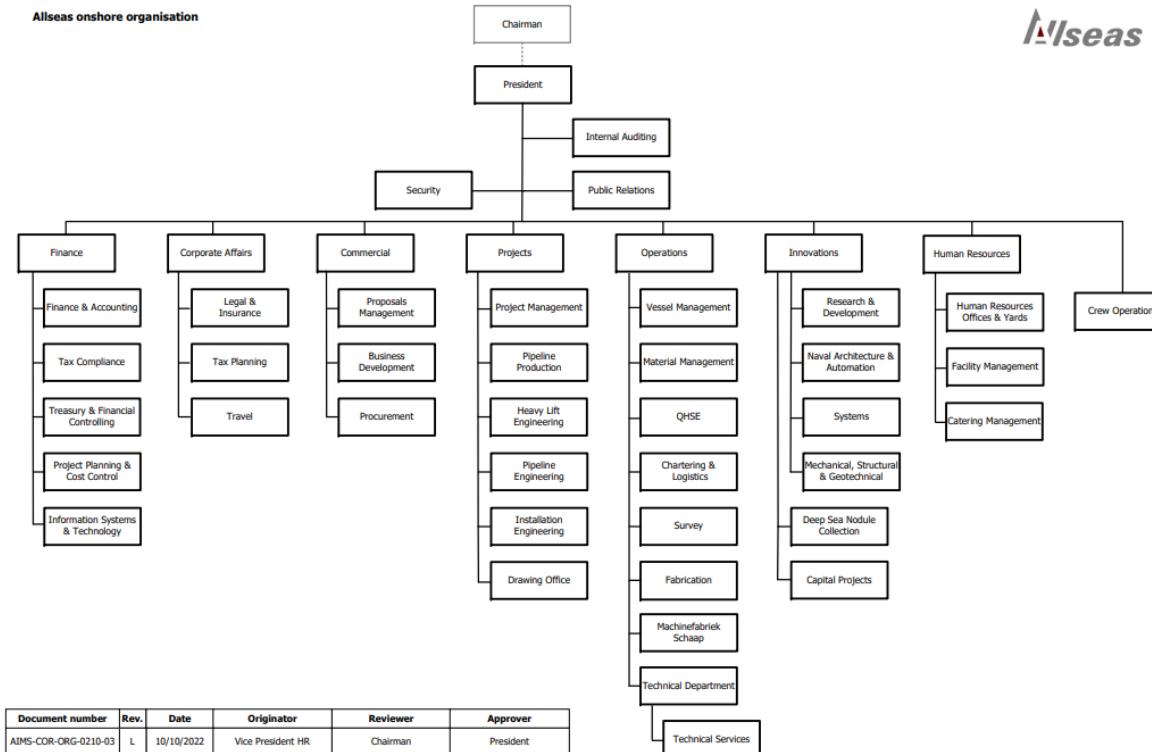


Figure 1: Allseas organizational structure.

The company's innovations department is responsible for the design and development of new technologies and solutions to improve the efficiency and safety of offshore operations. It is in this department that the internship project is situated. Figure 2 shows the organizational structure of the innovations department. In particular, the project is performed in the sub-department of Research & Development Delft, which is led by the department manager (Dominik Bujakiewicz). The project is supervised by the senior R&D engineer (J. Ramlakhan) and performed by the intern (Philip Soliman).

Organisation chart Innovations division

September 2024

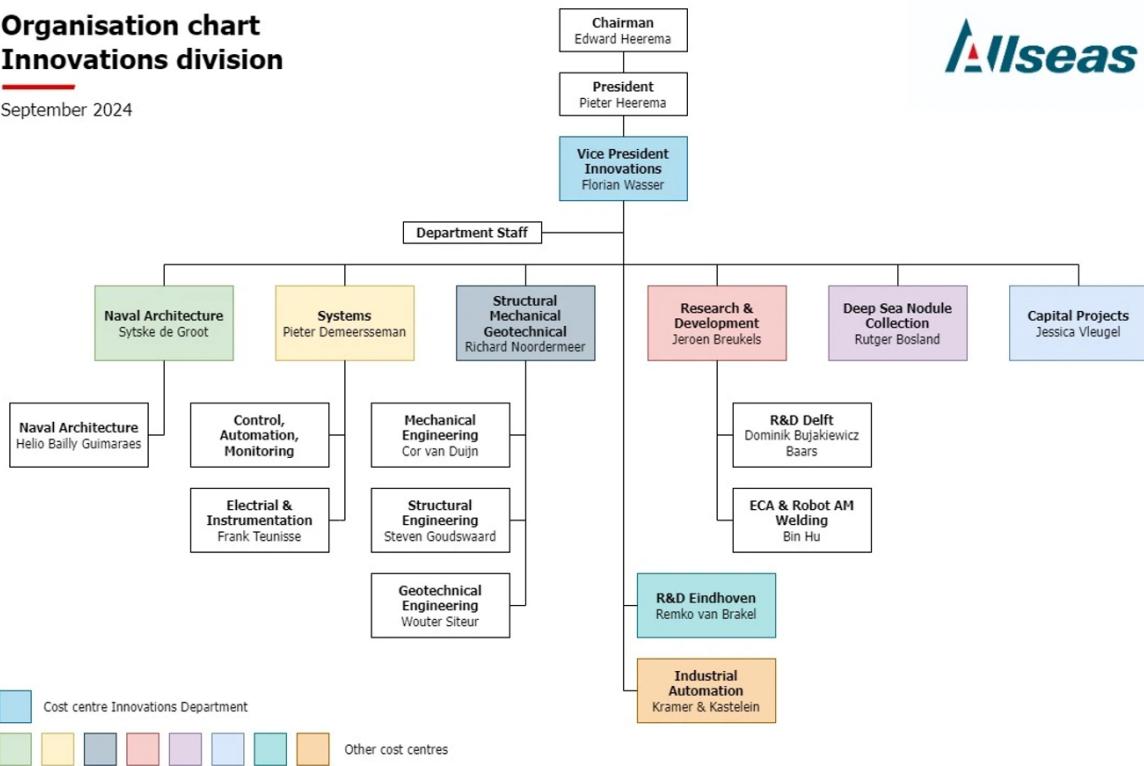


Figure 2: Allseas innovations department organizational structure.

2 Problem description

During pipelay every few minutes a new pipe (new joint) section must be fitted onto the end of the existing pipeline (main line). For decades in the offshore industry this has been a manual process. Allseas is working on a method to automate this.

2.1 Main requirements for automated line-up

Several projects have been initiated to automate the line-up process starting as early as 2008. However, These have all been unsuccessful. The latest project was started in 2017 and is still ongoing. The main requirements are:

1. Reducing line-up cycle times in the beadstall; This will result in significant cost reduction by decreasing vessel production time.
2. Increasing accuracy of line-ups (reduce amount of root pass weld errors); This will also reduce costs due to decreasing vessel production time.
3. Providing a safer work environment for the people working in the beadstall.

2.2 Current situation

The latest iteration of automated line-up is a system that uses 5 laser line scanners to measure the pipe ends' position and orientation (see figure 3).

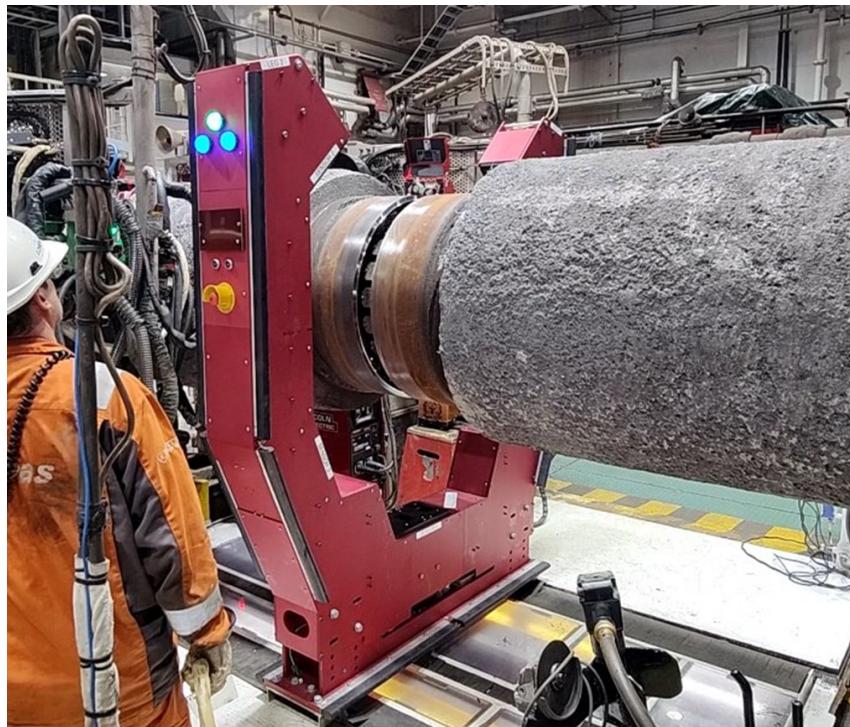


Figure 3: The portal or 'U-Frame'. The mainline is held in place in the background, while the new joint is lined up using instructions from automated line-up. The white stickers on the side of the frame indicate the position of the laser line scanners. Three on the side from which the picture is taken and two on the other side of the pipe. The U-Frame is mounted on the interface frame, which in turn is connected to the subframe. The subframe has motorized rollers and linear guides that allow the U-Frame to move with the new joint. This way the U-Frame ensures optimal visibility of the pipe ends for the laser line scanners. Additionally, the U-Frame can be moved out of way of other equipment, personnel or operations when not in use.

The laser line scanners generate 4192 data points in 5 different planes with a frequency of 60 ~ 70 Hz. Initially, a line finding algorithm derived from a Hough transform [3] is used to recover approximately straight sections from the laser's projection on the pipe. This allows for the recognition of the pipe ends or 'bevels' from the data.

Subsequently, the found pipe ends of the 5 sensors are combined to form a 3D representation of the pipe face. This is done using a combination of Newton-Raphson optimization and a least squares fit.

2.3 Internship goals

The internship aims to improve the current automated line-up system by:

- a **Decreasing code execution time** by searching for and replacing slow or redundant operations;
- b **Writing maintainable code** by refactoring the current codebase;
- c **Increasing machine accuracy** by comparing different pipe fitting algorithms.

3 Approach

In this section the approach to the problem is described. The approach is divided into several steps. The steps are ordered in a way that makes sense for the development of the project. The steps are as follows:

3.1 Understand automated line-up

Before any improvements can be made to the system, it is important to understand how the system works. This includes understanding the hardware, the software and the mathematical model that is used to fit the pipe ends. Next to that, getting to know the engineers that work on the project and the test set-up at the yard is also important. All of this provides a wholistic view of the system, which in turn highlights the signifigance of the improvements described in Section 2.3.

 section sections [4.1](#) to [4.3](#)

3.2 Familiarize with the code base

Read the code starting form the lowest level functions reading in the data received rom the laser line scanners to the highest level functions that combine the processed data to fit the pipe ends.

 section [4.4](#)

3.3 Create a profiler for the code

The best way to improve code speed is to know where the bottlenecks are. A profiler will help to identify these bottlenecks. Focusing on the most time consuming parts of the code prevents wasting resources on speeding up parts of the code that are not used often.

 internship goal [2.3 a](#)
 section [5](#)

3.4 Optimize the code

Once the major bottlenecks are identified, the next step is to find more efficient ways to perform the operations in question. Adjustments may involve reducing memory allocation, re-ordering operations, correcting or updating algorithms, implenting precomputation of constant values or optimising data locality.

 internship goal [2.3 a](#)
 section [6](#)

3.5 Identify unreadable code

Halfway during the development between 2017 and the time of this internship a new standard for clean, readable code has been introduced within Allseas. Parts of the old code have also been improved, but quite some code still needs a bit of restructuring. This step serves to identify the pieces of code that are least readable and need to be updated to the new Allseas coding standard.

 internship goal [2.3 b](#)
 sections sections [6](#) and [7](#)

3.6 Determine and improve upon current accuracy

Run new fitting algorithms on generated data to test the accuracy of the new algorithms. This will help to determine if the new algorithms are more accurate than the old ones. Do the same with recorded data from the test set-up to see if the new algorithms are also more reliable. Finally, test the algorithms on the real-time data from the test set-up to see if the new algorithms are also faster.

internship goal 2.3 c

section 7

3.7 Expand the fitting model to recognise ILUC

During line up of the new join with the main line, there is more equipment present in and around the latter. One major component present in the main line is the 'Internal Line Up Clamp' (ILUC) and is shown in Figure .

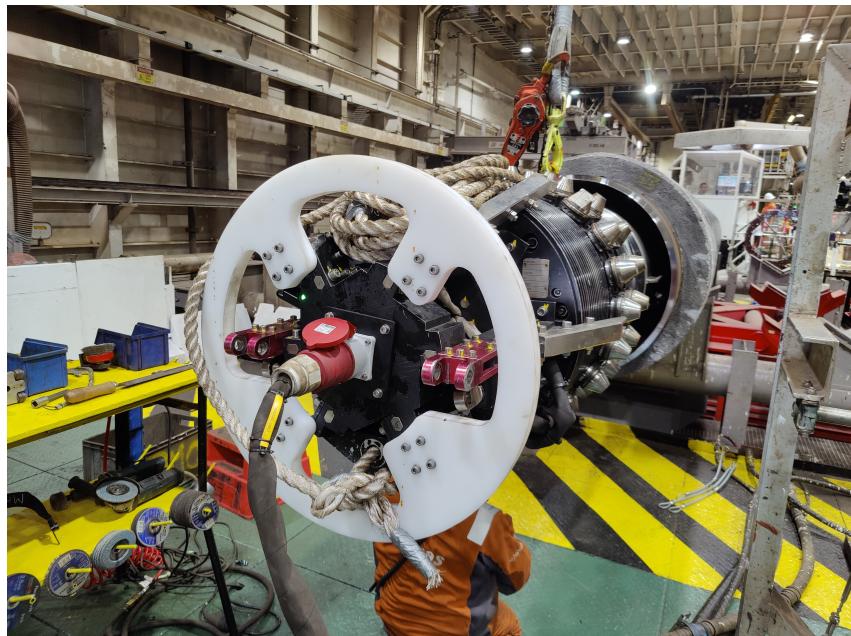


Figure 4: The Internal Line Up Clamp (ILUC). The ILUC sits in the main line, while the new join is shifted over it during line up. Its purpose is to steady the main line and new joint pipe ends after line up and during welding. Using a set of pneumatic knobs protruding from its exterior it pushes both pipe ends outward, ensuring a circular fit. Part of the ILUC is a black ribbed cylinder.

Recognising the ILUC gives a good indication of where the main line is before it can be measured, seeing as the U-Frame passes over the ILUC before it reaches the main line. Early estimation of the position and orientation of the main line benefits the overall speed of the process, because less corrections have to be made in the final line up stage.

Therefore, the final step of the internship is the development of a new part of the recognition software finds the a specific part of the ILUC; its ribbed cylinder. As before, accuracy and reliability can be tested using the test set-up in the yard. Furthermore, a combination of different kinds of line, plane and circle fitting algorithms based on minimization of geometric distances can be used for this. Lastly, Where pipe ends tend to have deformations due to manufacturing, pre-heating or bevelling, the ILUC is a near-perfect cylinder with a defined radius. This can be used to improve the accuracy of the fitting algorithms.

internship goal 2.3 c

section 8

4 Understanding line up automation

This section describes basic concepts of the line up automation (LUA) system. The automated line-up system is a system that aligns two pipes together. The system consists of a U-Frame, laser line scanners, Line up Car Assemblies (LCAs) and a sub-frame. The U-Frame houses the laser line scanners and is mounted on the sub-frame. The sub-frame is responsible for the U-Frame's movements. The LCAs are devices that control the new joint's position and orientation. The laser line scanners are sensors that detect the new joint's position.

4.1 The U-frame

Figure 5 shows the U-Frame. The U-frame is a part of the automated line-up system. It is precisely machined to ensure that the positions of the laser line scanners are known as accurately as possible.



Figure 5: The U-Frame. Most importantly, the U-frame houses five laser line scanners that are used to detect pipes and other objects like the ILUC. The U-Frame is mounted on an interface frame, which in turn is mounted on the sub-frame (not visible here). The sub-frame is responsible for the U-Frame's movements. To do ►change laser light to blue◀

4.2 Laser line scanners

The laserline scanners measure in 2 dimensions, width (x -direction) and depth (z -direction). The scanners each generate 4192 data points in any one measurement. The frequency of the measurements is only limited by how fast the laser line scanner runs (see section 4.4). The maximum scanning frequency of the lasers lies anywhere between 600 and 9000 Hz [7].

An important aspect of the laser line scanners is their Field of View (FOV). The FOV is the area in which the laser line scanners can detect objects. Figure 6 gives the FOV of the laser line scanners.

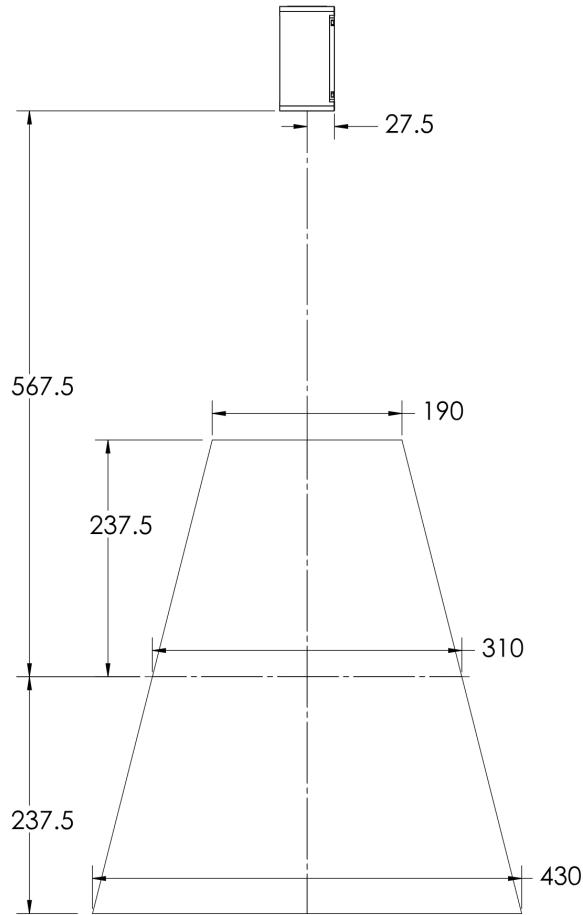


Figure 6: Field of View (FOV) of the laser line scanners. Measurements are in millimeters [8].

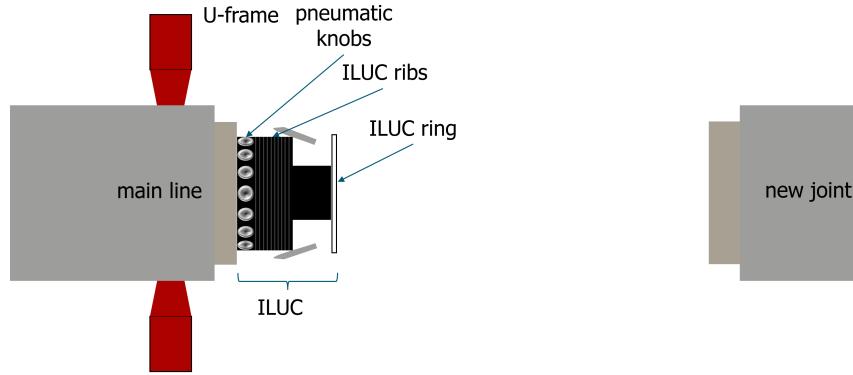
From figure 6 it can be seen that the FOV of the laser line scanners is limited. In particular, the width of the frame itself is ~ 20 cm, which means the laser line scanners can see ~ 10 cm beyond either side of the frame. This is why it is necessary for the U-Frame to move.

Depending on the distance of an object to a laser line scanner, the accuracy its measurement can vary. For an object close to the laser line scanners the resolution is $47\mu m$ in the x -direction. For an object further away from the laser line scanners the resolution is $104\mu m$ in the x -direction [7].

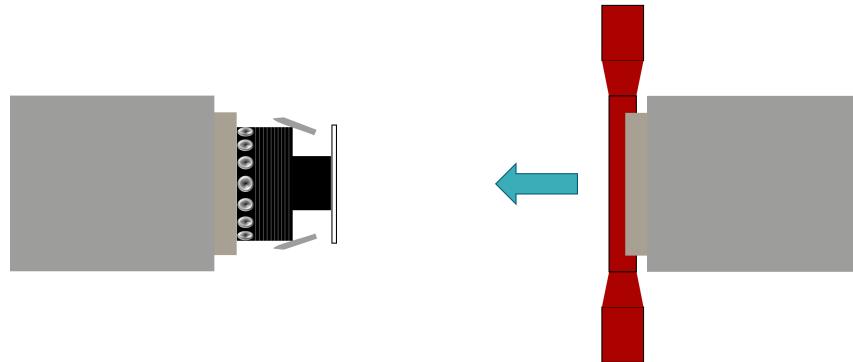
In terms of the z -direction, an indication of accuracy is given by the repeatability. This value is determined by using a flat target and calculating the 95% confidence variation of the average height over 4096 measurement frames. z values are averaged over the full FOV. The repeatability comes out to be $2\mu m$ [7].

4.3 Order of operations automated line-up

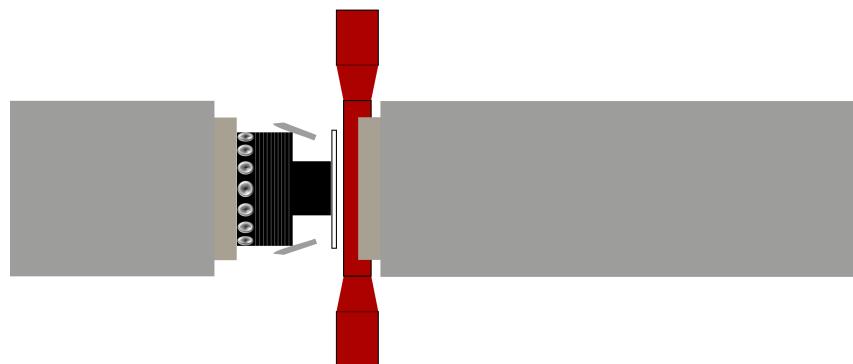
Figure 7 shows the order of operations during automated line-up.



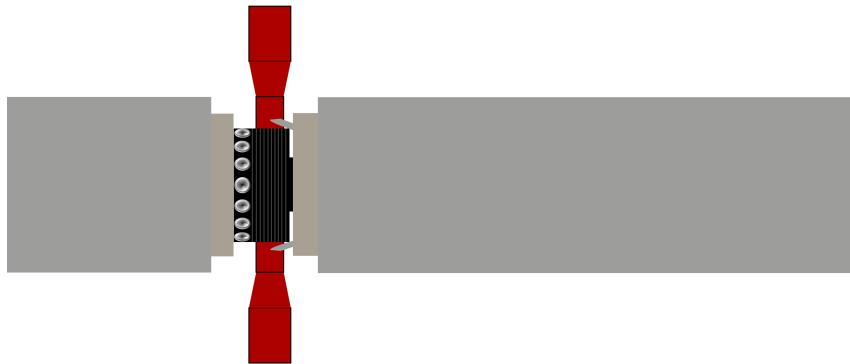
(a) Initial state of automated line up. The main line is depicted on the left, while the new joint is depicted on the right. The U-Frame is in its starting position. The ILUC sits in the main line. Visible parts of the ILUC are its pneumatic knobs, black ribbed cylinder (ILUC ribs) and white ring (ILUC ring).



(b) The U-Frame moves to the right until it encounters the new joint, which is moved towards the main line. The U-Frame reverses direction and tracks the new joint upon its recognition by the laser line scanners.



(c) The laser line scanners detect the ILUC ring. The first adjustments are made to the new joint's position. This is done to ensure that the new joint is shifted over the ILUC ring and other parts of the ILUC.



(d) The laser line scanners detect the ILUC ribs. This can be used to further adjust the new joint's position, matching more closely the orientation of the main line.



(e) Both pipe ends are fitted together during final line up. Adjustments are made to the new joint's position until a certain minimum gap tolerance is reached.



(f) U-Frame moves back to its starting position. The new joint is now fitted to the main line. The pipelines are now ready for welding.

Figure 7: Simplified order of operations during automated line-up. This procedure is repeated for each new joint that is added to the main line. After final line up in [7e](#) is completed, the new joint is welded to the main line. Then, the whole pipeline is moved forward and the process is repeated for the next new joint.

Figure 7 offers only simplified overview of the automated line-up process. Not shown in the figure are the;

- **sub-frame:** system that controls the U-Frame's movements;
- **laser line scanners:** sensors that detect the new joint's position;
- **Line up Car Assemblies (LCAs):** devices that precisely control the new joint's position and orientation;

4.4 Flow of line up automation software

The code flow of the laser line scanner software is shown in figure 8. This is an high-level overview of the software's operation.

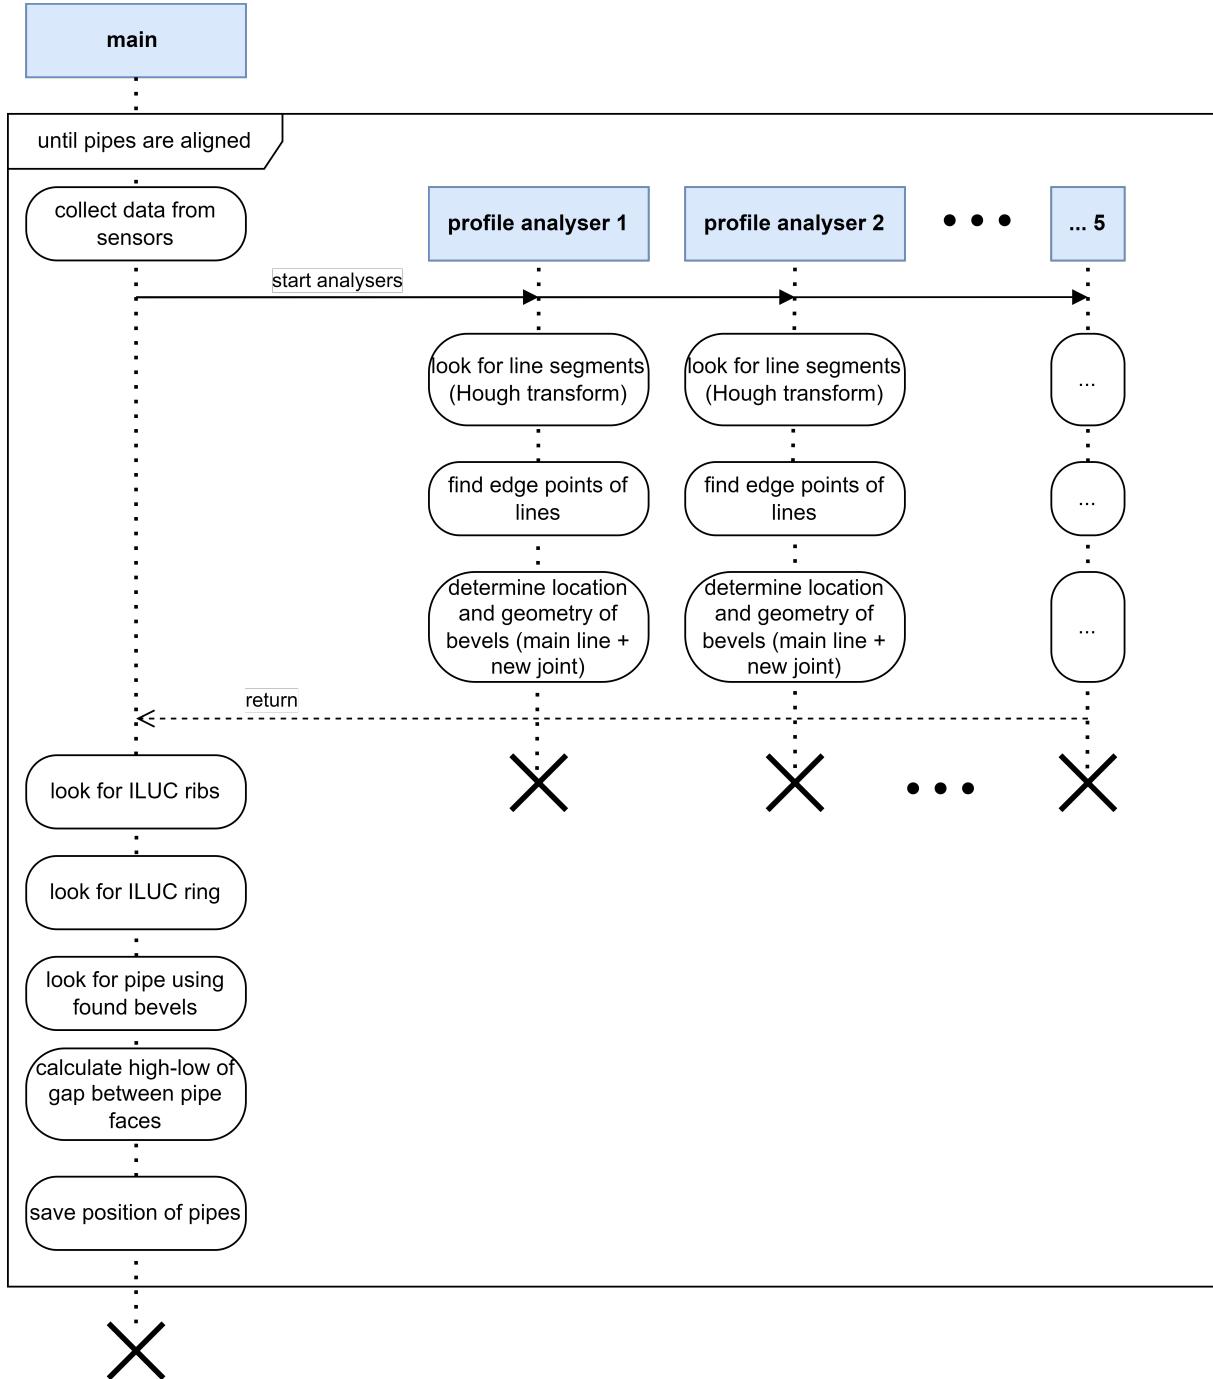


Figure 8: Flow of the laser line scanner software. The software is a multithreaded application. The blue boxes with bold text indicate the birth of these threads. The dotted lines are the 'life-lines' of the threads, which always terminate in an 'x'. The frame symbolises the main loop of the software. The top-left compartment in the frame specifies the loop-condition; 'until pipes are aligned'. The rounded boxes on the life-lines indicate functions that are called. Note that only the main thread can call functions. The main thread is responsible for the initialisation of the software and the creation of the other threads.

As can be seen in figure 8, the laser line scanner software assigns a separate thread for the analysis of the data coming from each of the five laser line scanners. These analyser

threads are responsible for finding line-segments, edge points and pipe bevels in their respective 2-dimensional data sets. Once this preliminary, disjoint analysis is done, the main thread combines the processed data into a 3-dimensional model of the ILUC ribs, -ring and/or pipes. Specifically, the models for the pipes are used to calculate the high-low values of the gap between the new joint and the main line. The high-low values determine if the pipes are lined-up correctly. If the gap is within a certain tolerance, the software will signal that the pipes are aligned and it terminates the main loop.

Comparing Figures 7 and 8, it becomes clear that the relative activity of the various parts of the software change depending on the stage of the automated line-up process. For example, determining the location and geometry of a pipe's bevel is only necessary as soon as the new joint comes into Field of View (FOV) of the laser line scanners (step 7b). Then, the activity roughly doubles as soon as the main line comes into FOV of the laser line scanners during final line up (step 7e). Furthermore, the finding of lines and their edges in each of the profile analyser threads becomes more computationally expensive for the ILUC ribs. This is further explored in section 5.

5 Profiler

This section describes the use of a profiler to identify bottlenecks in the LUA software. The profiler is used to identify functions that consume a lot of CPU time, and thus are candidates for optimization.

Microsoft Visual Studio (VS) comes with an installation of VSDiagnostics; a command-line profiling tool [6]. VSDiagnostics generates *.diagsession files, which VS interprets and visualizes. For this reason, it is more convenient to call VSDiagnostics indirectly from within Visual Studio.

To do this, debug and symbol flags need to be added to the compilation commands of object files and final linking of the main executable. Adding these flags ensures the compiler generates symbol (for Windows .pdb) files, which are required by any profiler.

The VS profiler discerns between two categories of CPU time: self-time and total time. Self-time is the time spent in a function, excluding time spent in subroutines/-functions. Total time is the time spent in a function, including time spent in subroutines/-functions [5]. The self time is expressed as

$$\text{self-time} = \frac{\text{CPU time spent in the function}}{\text{total execution time}} \times 100\%. \quad (1)$$

5.1 Profiling test setup

Profiling measurements are done in a controlled environment, where the LUA software is run with one laser line scanner. The scanner is placed in front of a test setup consisting of a lone ILUC ring, a mockup for the ILUC ribs and new joint.



(a) ILUC ring.



(b) ILUC ribs.



(c) Full setup. Note, laser is not connected

Figure 9: The experimental setup for profiling measurements.

These objects are moved in a controlled manner to simulate the phased of LUA. The experimental setup is described in Table 1.

Phase	Description	Experimental Setup
7b Joint recognition	LUA software detects joint as it moves towards mainline. U-Frame moves with joint after (successful) detection.	Flat surface resembling the joint (back of the ILUC-ribs mockup) is quickly brought into frame and subsequently kept still.
7c ILUC ring recognition	Joint and U-Frame move until ILUC ring is detected, after which partial line up takes place.	Next to the joint mockup, the side of the ILUC ring is brought into frame representing the ILUC-ring.
7d ILUC ribs recognition	The ILUC-ring moves into the joint during the aforementioned partial line-up, after which the ILUC-ribs are detected. The joint is moved over the ribs towards the mainline.	A joint mockup (a side of the ILUC-ring mockup) is still in frame. The ILUC ribs mockup is slowly moved into frame, revealing all its ribs.
7e Final line up	The mainline gets detected and the joint is lined up with it	The joint mockup and a second flat surface (mainline mockup) are slowly moved together.

Table 1: Profiling experiments simulating essentials phases of LUA. Each of the experiments are performed with one laser line scanner and done on a standard company PC (HP Z4 G4).

Every profiling test is approximately 20 seconds in length, in keeping with the profiling guidelines. Every test is preceded by a 20 second settling period during which the LUA software runs, but no profiling or setup movement takes place. The profiling sampling rate is ≥ 1000 Hz.

5.2 Profiling results

Shown in Table 2 are the most time consuming functions per LUA phase.

Phase	Top function	self-time (%)
Joint recognition	<code>LineFinder::calculateNrPointsInEachBin</code>	13.94
ILUC ring recognition	<code>LineFinder::calculateNrPointsInEachBin</code>	17.89
ILUC ribs recognition	<code>LineFinder::calculateNrPointsInEachBin</code>	18.53
Final line up	<code>LineFinder::calculateNrPointsInEachBin</code>	15.87

Table 2: Profiling results of experiments described in Table 1.

The results show that the function `LineFinder::calculateNrPointsInEachBin` is the most time consuming function in all phases. This function is a bottleneck and is a target for optimization.

As mentioned in section 4.4, recognition of the ILUC ribs is the most time consuming phase. This can be attributed to the ILUC ribs consisting of many smaller line segments. The Hough transform function executed in each of the profile analyser threads in figure 8 is ran for each of these line segments. This in turn results in many calls to the function `LineFinder::calculateNrPointsInEachBin`, which is a part of the Hough transform.

Figure 10 shows profiling results of parts of the bottleneck function

```

22 (0.07%) 77 const float cosThisAngle = cos( thisAngle );
5 (0.02%) 78 const float sinThisAngle = sin( thisAngle );
79
80 // Allocate the arrays within this function for faster access
231 (0.71%) 81 const auto dataPointX = m_profileData->m_profileDataX;
299 (0.92%) 82 const auto dataPointZ = m_profileData->m_profileDataZ;
5 (0.02%) 83 const int numberofValidPoints = m_profileData->m_nrValidDataPoints;
319 (0.99%) 84 const auto profilePointID = m_profileData->m_profilePointID;
85
717 (2.22%) 86 for ( int xIndex = 0; xIndex < numberofValidPoints; xIndex++ )
87 {
88     // Check if the point is not already assigned to a line.
2030 (6.28%) 89 if ( profilePointID[xIndex] == 0 )
90 {
91     {
92         // Check the distance from the point on the search line to the origin.
1634 (5.05%) 93 float thisDistance = cosThisAngle * dataPointX[xIndex] + sinThisAngle * dataPointZ[xIndex];
94
95         // Check if the distance is within the search range.
292 (0.90%) 96 if ( thisDistance >= START_SEARCH_DISTANCE && thisDistance < END_SEARCH_DISTANCE )
97 {
98             // Calculate in which bin the point belongs and add it to the correct bin.
343 (1.06%) 99 int binSelector = static_cast<int>( ( thisDistance - START_SEARCH_DISTANCE ) / SEARCH_DIST_BIN_SIZE
1209 (3.74%) 100 nrPointsThisDistance[binSelector] += 1;
101
102     }
103 }
104 }

```

Figure 10: Profiling results of the function `LineFinder::calculateNrPointsInEachBin` during the ILUC ribs recognition phase.

From figure 10 it becomes clear what the most time consuming operations are in the calculation of the Hough transform. In order of most time consuming to least, these are:

- B.N. 1 Checking if a data point is not already associated to a line ([line 89 in figure 10](#));
- B.N. 2 Calculating the distance between a data point and the laser line scanner's origin ([line 93 in figure 10](#));
- B.N. 3 Determining in which distance bin a data point falls ([line 100 in figure 10](#)).

6 Optimize & refactor code

In section 5.2 it is shown that the function `LineFinder::calculateNrPointsInEachBin` is the most time consuming function in all phases of the LUA software.

Before optimization can take place it is essential to understand the structure of this part of the LUA software. Algorithm 1 shows the pseudo code of the function `LineFinder::calculateNrPointsInEachBin`.

Algorithm 1 Pseudo code of the function `LineFinder::calculateNrPointsInEachBin`.

```
1: function CALCULATENRPOINTSINEACHBIN( $\alpha$ )
2:   calculate sine and cosine of this angle  $\alpha$ 
3:   fetch X Z coordinates of the data points
4:   fetch the indices of the data points (used to determine line association)
5:   for each data point do
6:     check if data point is not already associated to a line
7:     calculate distance between data point and laser line scanner's origin
8:     determine in which distance bin a data point falls
9:   end for
10: end function
```

Algorithm 1 is nested in a loop that iterates over search angles specified for the Hough transform. This is shown in algorithm 2

Algorithm 2 Pseudo code of the function `LineFinder::calculateBinContainingMostPoints`.

```
1: function CALCULATEBINCONTAININGMOSTPOINTS
2:   for each angle  $\alpha$  in all search angles do
3:     calculateNrPointsInEachBin( $\alpha$ )
4:     for each distance bin in all search bins do
5:       if bin contains more points than the current maximum then
6:         update the current maximum
7:       end if
8:     end for
9:   end for
10: end function
```

With regards to the order of for-loops, it is optimal to have a larger inner for-loop and a smaller outer for-loop [4, lecture 7, "temporal locality"]. Algorithms 1 and 2 show the nested for-loop structure that is used to perform the Hough transform. Note that the number of search angles is approximately 30, while the number of data points is approximately 4000. This means that the inner for-loop is much larger than the outer for-loop. Hence, the order of the for-loops is already optimal.

Closer inspection of algorithm 1 reveals that the fetching operations in lines 2 and 3 do not depend on the angle α . This means that these fetching operations can be moved outside the for-loop to the beginning of algorithm 2. This ultimately results in less memory calls happening in the inner for-loop. Furthermore, the calculation of the sine and cosine can be done once for all angles in the main thread of the LUA software before the main program loop (figure 8).

Figure 11 shows the profiling results of the implemented prefetching and precalculation optimizations.

```

112 {
9 (0.02%) 113     const int numberOfValidPoints = m_profileData->m_nrValidDataPoints;
114
115     // Create bins for all distances to keep track of how many datapoints fall into each interval
116     std::array<unsigned, NUM_ROUGH_RADII> nrPointsThisDistance = { 0 };
117
118     for ( int xIndex = 0; xIndex < numberOfValidPoints; xIndex++ )
119     {
120         // Check if the point is not already assigned to a line.
121         if ( profilePointID[xIndex] == 0 ) // line association check
122         {
123             {
124                 // Check the distance from the point on the search line to the origin.
125                 float thisDistance = cosine * dataPointX[xIndex] + sine * dataPointZ[xIndex]; // distance calculation
126
127                 // Check if the distance is within the search range.
128                 if ( thisDistance >= START_SEARCH_DISTANCE && thisDistance < END_SEARCH_DISTANCE )
129                 {
130                     // Calculate in which bin the point belongs and add it to the correct bin.
131                     int binSelector = static_cast<int>( ( thisDistance - START_SEARCH_DISTANCE )
132                                                 / SEARCH_DIST_BIN_SIZE ); // bin selection
133                     nrPointsThisDistance[binSelector] += 1; // bin association
134                 }
135             }
136         }
137     }
138
139     return nrPointsThisDistance;

```

(a) Profiling results of the function LineFinder::calculateNrPointsInEachBin (algorithm 1) after precalculation optimization.

```

51     // prefetch arrays & vectors before going into loop over angles
52     const std::array<float, NUM_ROUGH_ANGLES> angles = m_analyserSettings->roughAngles;
53     const std::array<float, NUM_ROUGH_ANGLES> cosines = m_analyserSettings->cosines;
54     const std::array<float, NUM_ROUGH_ANGLES> sines = m_analyserSettings->sines;
55
56     const std::vector<float> dataPointX( m_profileData->m_profileDataX.cbegin(), m_profileData->m_profileDataX.cend() );
57     const std::vector<float> dataPointZ( m_profileData->m_profileDataZ.cbegin(), m_profileData->m_profileDataZ.cend() );
58
59
60     BestRoughLineFit bestFit;
61
62     for ( unsigned angleNr = 0; angleNr < NUM_ROUGH_ANGLES; angleNr++ )
63     {
64         auto tempBestFit = calculateDistanceIntervalContainingMostPoints(
65             angles[angleNr], cosines[angleNr], sines[angleNr], dataPointX, dataPointZ, profilePointID );
66
67         // Use the angle with the first best result
68         if ( tempBestFit.nrOfPointsFromCombinedBins > bestFit.nrOfPointsFromCombinedBins )
69         {
70             bestFit = tempBestFit;
71         }
72     }
73
74     return bestFit;
75
76

```

(b) Profiling results of the function LineFinder::calculateBinContainingMostPoints (algorithm 2) after prefetching optimization.

Figure 11: Profiling results of optimized bottleneck functions. Calculation of the sinusoidal functions is done during initialisation of the LUA software. Data arrays are prefetched before before the inner for-loop.

Comparing lines 55-57 from figure 11b with lines 81-84 from figure 10 shows that both the percentage of CPU time and the number of times these lines are sampled by the profiler are lower in the optimized case. This is a clear indication that the optimizations have been successful.

Further optimizations can be made by avoiding the most time consuming operation in the inner for-loop **Bottle Neck 1**. In general an if-statement in a for-loop causes a ‘pipeline stall’ [4, lecture 7, “Pipeline stall”, slide 19]. This is because the CPU must wait until the if-statement is evaluated before it can continue executing the next instruction. This effect is made worse by the fact that the if-statement is in the inner most nested for-loop. The solution to this problem is to use a data structure that allows for fast look-up of data points that are already associated to a line. Perhaps this can be done by using a hash table. Alternatively, algorithmic adjustments can be made to avoid the if-statement entirely. In any case, this is left as a topic for future work.

7 Accuracy of circle fits

This section gives a mathematical estimate for the accuracy of circle fit algorithms that minimize geometric distances. This is useful knowlegde for LUA, since this kind of ‘geometric circle fit’ algorithm is used to form a 3-dimensional representation of the pipes, among other things (section 4.4). This discussion then ends with the implementation in the LUA software of a new kind of circle fit algrorithm that has a variance that is the same as the least squares method, but no bias in its estimation of the circle radius. This section builds on the methods described in the paper by Al-sharadqah and Chernov [2].

7.1 Statistical model

In order to estimate the accuracy of a circle fit, we need to define a statistical model for the error in the input. To that end, we assume that the data points (x_i, y_i) are noisy observations of the true points $(\tilde{x}_i, \tilde{y}_i)$, where the noise is i.i.d. normally distributed with bias 0 and variance s^2 [2, section 2].

$$x_i = \tilde{x}_i + \delta x_i, \quad y_i = \tilde{y}_i + \delta y_i, \quad \delta x_i, \delta y_i \sim N(0, s^2). \quad (2)$$

This gives that the true points must lie on a circle, satisfying the equation

$$(\tilde{x}_i - \tilde{a})^2 + (\tilde{y}_i - \tilde{b})^2 = \tilde{R}^2 \quad \text{for } i = 1, \dots, n, \text{ midpoint } (\tilde{a}, \tilde{b}) \text{ and radius } \tilde{R}. \quad (3)$$

In the case of LUA the angles of the laser line scanners in the U-frame are known (section 4.2). This suggests normalisation of the data points to the unit circle, which is done by

$$u_i = \frac{\tilde{x}_i - \tilde{a}}{\tilde{R}} = \cos \phi_i, \quad v_i = \frac{\tilde{y}_i - \tilde{b}}{\tilde{R}} = \sin \phi_i, \quad (4)$$

where ϕ_i is the angle corresponding to the true point $(\tilde{x}_i, \tilde{y}_i)$ projected on the pipe (or ILUC ring/ -ribs) by the i^{th} laser line scanner. This leads to the normalised data matrix W with rows $W_i = (u_i, v_i, 1)^T$.

Define the vector of estimates $Q = (a, b, r)^T$ for the circle parameters. Then, the first order terms in the Taylor expansion of Q in terms of $\delta x_i, \delta y_i$ is denoted by $\Delta_1 Q$. The accuracy of the circle fit can then be estimated by the variance of $\Delta_1 Q$ [2, section 2].

7.2 Accuracy estimation

The geometric circle fit has both bias and (co-)variance. under the assumption that the number of data points approaches infinity $n \rightarrow \infty$, the bias is given by

$$\mathbb{E}[\Delta_1 Q] = \frac{2s^2}{\tilde{R}} \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix} \quad [2, \text{ equation 6.14}], \quad (5)$$

and therefore the bias exists only in the radius estimate as $\mathbb{E}[\Delta_1 R] = \frac{2s^2}{\tilde{R}}$. The (co-)variance is given as

$$\text{cov}(\Delta_1 Q) = s^2 (W^T W)^{-1} \quad [2, \text{ equation 8.12}]. \quad (6)$$

Taking the diagonal elements of the covariance matrix gives the variance of the estimates for the circle parameters. Taking the square root of these variances gives the standard deviation $s(\cdot)$ of the estimates. The latter is good indicator of the accuracy of the circle fit.

It remains to estimate the noise level s in the data and the true angles ϕ_i for all i . From section 4.2 it is known that the resolution of the laser line scanners is $47\mu\text{m}$ to $104\mu\text{m}$ in the x -direction and 2μ in the z -direction. This suggests that the noise level s is on the order of 10^{-4}

m. This is a rough estimate, but it is sufficient for the purposes of this discussion. As stated in the previous section 7.1, the true angles correspond to the angle of the laser line scanner in the U-frame.

Using the above circle fit accuracy can be calculated for the idealised scenario where a 24-inch (= 0.5 m) pipe is exactly in the center of the U-frame with zero pitch or yaw. In this case

$$\begin{aligned}\tilde{a} &= \tilde{b} = 0, \quad \tilde{R} = 0.5 [m], \\ \phi_1 &= -\frac{\pi}{4}, \\ \phi_2 &= \frac{\pi}{4}, \\ \phi_3 &= \frac{5\pi}{8}, \\ \phi_4 &= -\frac{5\pi}{8}, \\ \phi_5 &= \pi,\end{aligned}$$

where the angles are measured from the z -axis in the y - z plane with positive rotation defined by the right hand rule applied to the x -axis. In this simple case the angles are the same as the angles of the laser line scanners in the U-frame.

Substituting this into the formulae eqs. (5) and (6) above gives

$$\begin{aligned}E[\Delta_1 R] &= , \\ s(a) &= \sqrt{(\text{cov}(\Delta_1 Q))_{11}} = , \\ \text{cov}(\Delta_1 b) &= \sqrt{(\text{cov}(\Delta_1 Q))_{22}} = , \\ \text{cov}(\Delta_1 R) &= \sqrt{(\text{cov}(\Delta_1 Q))_{33}} = .\end{aligned}$$

8 ILUC Recognition

9 Conclusion

Perform tests in the yard with the complete system, to test new code for points described in proposed approach: - Create a profiler for the code to search for slow operations. Milestone: profiler gives insight into performance of different parts of the code and allows comparison between different versions of said code. This way we can determine which is optimal - Find more efficient ways to perform the needed operations. Milestone: Increase overall code speed. Use profiler as quantification - Isolate parts of the mathematical model and test accuracy. Milestone: Mathematical as well as experimental estimation for line-up accuracy meets 0.1 mm requirement - Expand the fitting model to increase the workability of the complete system. Milestone: Successful and efficient recognition of ribbed cylinder allows for good pre-line-up of pipes bsectionProfiler

10 Skills learned

Bibliography

- Allseas. (2024). *About allseas*. [URL](#).
- Al-sharadqah, A., & Chernov, N. (2009). Error analysis for circle fitting algorithms. *Electronic Journal of Statistics*, 3. <https://doi.org/10.1214/09-EJS419>
- Duda, R. O., & Hart, P. E. (1972). Use of the hough transformation to detect lines and curves in pictures. *Commun. ACM*, 15(1), 11.15. <https://doi.org/10.1145/361237.361242>
- Lin, D. I. H. (2024, April). Lecture notes in scientific programming (tw4260tu).
- microsoft. (2024a). *Profiling cpu usage*. [URL](#).
- microsoft. (2024b). *Vs diagnostics*. [URL](#).
- Technologies, L. (2024a). *Gocator 2650 smart 3d laser line profile sensor*. [URL](#).
- Technologies, L. (2024b). *Gocator 2650 smart 3d laser line profile sensor*. [URL](#).